

# Linear Regression

Quách Đình Hoàng

# Outline

- Empirical Risk Minimization
- Regression
- Linear Regression
- Regularized Regression
- Hyperparameter Tuning

# Empirical Risk Minimization – ERM

- ERM is a common framework for supervised learning
- Given:
  - some labelled training dataset  $D = \{(x^{(i)}, y^{(i)})\}_{i=1}^n$
  - a loss function  $l: Y \times Y \rightarrow R$
  - a hypothesis class or set of functions  $F$
- The goal is to find

$$\hat{f} = \underset{f \in F}{\operatorname{argmin}} \sum_{i=1}^n l(f(x^{(i)}), y^{(i)})$$

with the hope that

$$E_{p(x,y)}[l(f(x), y)] \approx \frac{1}{n} \sum_{i=1}^n l(f(x^{(i)}), y^{(i)})$$

# Empirical Risk Minimization – ERM

- ERM is a common framework for supervised learning

- Given:

- a labelled training dataset  $D = \{(x^{(i)}, y^{(i)})\}_{i=1}^n, x^{(i)} \in \mathbb{R}^{d+1}$
- a loss function  $l: Y \times Y \rightarrow R$
- a hypothesis class or set of functions  $F$

- The goal is to find

$$\hat{f} = \underset{f \in F}{\operatorname{argmin}} \sum_{i=1}^n l(f(x^{(i)}), y^{(i)})$$

- Depending on the choice of  $F$  and  $l$ , this objective function may be **convex** (easy to optimize) or **non-convex** (hard)

# Regression

- Regression is a type of supervised learning
- Given:
  - a labelled training dataset  $D = \{(x^{(i)}, y^{(i)})\}_{i=1}^n$
  - a loss function  $l: Y \times Y \rightarrow R$ , where  $Y = \mathbb{R}$
  - a hypothesis class or set of functions  $F$

- The goal is to find

$$\hat{f} = \underset{f \in F}{\operatorname{argmin}} \sum_{i=1}^n l(f(x^{(i)}), y^{(i)})$$

- Depending on the choice of  $F$  and  $l$ , this objective function may be **convex** (easy to optimize) or **non-convex** (hard)

# Linear Regression (Ordinary Least Squares)

- Linear regression is a simplest type of regression

- Given:

- a labelled training dataset  $D = \{(x^{(i)}, y^{(i)})\}_{i=1}^n$

- a loss function  $l(y, y') = (y - y')^2$

- a hypothesis class or set of functions  $F$  of the form  $f(x) = w_0 + \sum_{i=1}^d w_i x_i$

- The goal is to find

$$\hat{f} = \underset{f \in F}{\operatorname{argmin}} \sum_{i=1}^n l(f(x^{(i)}), y^{(i)})$$

# Linear Regression (Ordinary Least Squares)

- Classification is a type of supervised learning

- Given:

- a labelled training dataset  $D = \{(x^{(i)}, y^{(i)})\}_{i=1}^n$
- a loss function  $l(y, y') = (y - y')^2$
- a hypothesis class or set of functions  $F$  of the form  $f(x) = w_0 + \sum_{i=1}^d w_i x_i = w^T x$

$$w^T = (w_0, w_1, \dots, w_d)$$
$$x^T = (1, x_1, \dots, x_d)$$

- The goal is to find

$$\hat{w} = \underset{w}{\operatorname{argmin}} \sum_{i=1}^n (w^T x^{(i)} - y^{(i)})^2$$

# Linear Regression (Ordinary Least Squares)

- Linear regression is a simplest type of regression
- Given:

- a labelled training dataset  $D = \{(x^{(i)}, y^{(i)})\}_{i=1}^n$
- a loss function  $l(y, y') = (y - y')^2$
- $F$  = all functions of the form  $f(x) = w_0 + \sum_{i=1}^d w_i x_i = w^T x$
- The goal is to find

$$\begin{aligned} w^T &= (w_0, w_1, \dots, w_d) \\ x^T &= (1, x_1, \dots, x_d) \end{aligned}$$

$$\hat{w} = \underset{w}{\operatorname{argmin}} \sum_{i=1}^n (w^T x^{(i)} - y^{(i)})^2$$

or

$$\hat{w} = \underset{w}{\operatorname{argmin}} \|Xw - y\|^2$$

where:

$$X = \begin{bmatrix} x^{(1)T} \\ \vdots \\ x^{(n)T} \end{bmatrix}, y = \begin{bmatrix} y^{(1)} \\ \vdots \\ y^{(n)} \end{bmatrix}, x^{(i)} = \begin{bmatrix} 1 \\ x_1^{(i)} \\ \vdots \\ x_d^{(i)} \end{bmatrix}, w = \begin{bmatrix} w_0 \\ w_1 \\ \vdots \\ w_d \end{bmatrix}$$



# Linear Regression (Ordinary Least Squares)

$$L_D(w) = \|Xw - y\|^2 = (Xw - y)^T(Xw - y)$$

⋮

$$\hat{w} = (X^T X)^{-1} X^T y$$

# Regularized Linear Regression

- Regularized empirical risk minimization that penalizes model complexity to deal with overfitting
- Given:
  - some labelled training dataset  $D = \{(x^{(i)}, y^{(i)})\}_{i=1}^n$
  - a loss function  $l: Y \times Y \rightarrow R$ , where  $Y = R$
  - a hypothesis class or set of functions  $F$
  - a regularizer function  $r: w \rightarrow R$
  - a regularization parameter  $\lambda$
- The goal is to find

$$\hat{w} = \underset{w}{\operatorname{argmin}} \sum_{i=1}^n l(f_w(x^{(i)}), y^{(i)}) + \lambda r(w)$$

# Regularized Linear Regression

- Linear regression

$$L_D(w) = \sum_{i=1}^n (w^T x^{(i)} - y^{(i)})^2$$

- Ridge regression ( $L_2$  regularization)

$$\underset{w}{\text{minimize}} \left( L_D(w) + \frac{\lambda}{2} \|w\|_2^2 \right)$$

$\lambda$  is hyperparameter

- Lasso regression ( $L_1$  regularization)

$$\underset{w}{\text{minimize}} \left( L_D(w) + \frac{\lambda}{2} \|w\|_1 \right)$$

$\lambda$  is hyperparameter

- ElasticNet regression ( $L_1 + L_2$  regularization)

$$\underset{w}{\text{minimize}} \left( L_D(w) + \lambda \left( \alpha \|w\|_1 + \frac{1-\alpha}{2} \|w\|_2^2 \right) \right)$$

$\lambda$  and  $\alpha$  are hyperparameters

# Ridge Regression

$$L_D(w) = \|Xw - y\|^2 + \lambda\|w\|^2 = (Xw - y)^T(Xw - y) + \lambda w^T w$$

$\vdots$

$$\hat{\theta} = (X^T X + \lambda I_{d+1})^{-1} X^T y$$

$I_{d+1}$ : is the  $(d + 1) \times (d + 1)$  identity matrix

# Gradient Descent

- We're trying to minimize some function  $L$
- Suppose at iteration  $t$ : we're get  $w^{(t)}$
- With learning rate  $\eta$ , we update  $w^{(t)}$  (at iteration  $t + 1$ ) as follow

$$w^{(t+1)} = w^{(t)} - \eta \nabla_w L(w^{(t)})$$

# Gradient Descent for Linear Regression

- Dataset  $D = \{(x^{(i)}, y^{(i)})\}_{i=1}^n$

1. Initialize  $w^{(0)} = 0$  (zero vector) and set  $t = 0$

2. While **not converged**

- Compute the gradient

$$\nabla_w L_D(\theta^{(t)}) = 2X^T X w^{(t)} - 2X^T y$$

- Update the weights

$$w^{(t+1)} = w^{(t)} - \eta \nabla_w L_D(w^{(t)})$$

- Increment  $t$ :

$$t = t + 1$$

- Output  $w^{(t)}$

# Gradient Descent for Linear Regression – Change Step Size

- Dataset  $D = \{(x^{(i)}, y^{(i)})\}_{i=1}^n$
- 1. Initialize  $w^{(0)} = 0$  (zero vector) and set  $t = 0$
- 2. While **not converged**
  - Compute the gradient

$$\nabla_w L_D(w^{(t)}) = 2X^T X w^{(t)} - 2X^T y$$

- Update the weights

$$w^{(t+1)} = w^{(t)} - \eta_t \nabla_w L_D(w^{(t)}) \quad \eta_t = \frac{\eta_0}{n\sqrt{t+1}}$$

- Increment  $t$ :

$$t = t + 1$$

- Output  $w^{(t)}$

# Gradient Descent for Linear Regression – Change Step Size

- Dataset  $D = \{(x^{(i)}, y^{(i)})\}_{i=1}^n$
- 1. Initialize  $w^{(0)} = 0$  (zero vector) and set  $t = 0$
- 2. While **not converged**
  - Compute the gradient

$$\nabla_w L_D(w^{(t)}) = 2 \sum_{i=1}^n (w^{(t)T} x^{(i)} - y^{(i)}) x^{(i)}$$

- Update the weights

$$w^{(t+1)} = w^{(t)} - \eta_t \sum_{i=1}^n (w^{(t)T} x^{(i)} - y^{(i)}) x^{(i)} \quad \eta_t = \frac{\eta_0}{n\sqrt{t+1}}$$

- Increment  $t$ :


$$t = t + 1$$

- Output  $w^{(t)}$



# Hyperparameter Tuning

- Suppose we want to compare multiple hyperparameter settings  $\lambda_1, \lambda_2, \dots, \lambda_k$
- For  $i = 1, 2, \dots, k$ 
  - Train a model on  $D_{train}$  using  $\lambda_i$
- Evaluate each model on  $D_{val}$  and find the best hyperparameter setting,  $\lambda_{i^*}$
- Compute the error of a model trained with  $\lambda_{i^*}$  on  $D_{test}$



# Summay

- Empirical Risk Minimization
- Regression
- Linear Regression
- Regularized Regression
- Hyperparameter Tuning